

Analysis of Min-Heap Implementation (Partner B)

Algorithm Overview

A binary *min-heap* is a complete binary tree in which every parent node has a key less than or equal to those of its children. The shape property ensures the tree is always balanced—except possibly at the last level, which is filled from left to right. The heap property enables efficient access to the smallest element at the root.

In an array-based representation, the children of the node at index i reside at $2*i + 1$ and $2*i + 2$, and the parent of i is at $\text{floor}((i - 1)/2)$. The core operations supported by the student's implementation are:

- **insert**: appends an element to the end of the array and performs *sift-up* (heapify-up) to restore the heap property.
- **extractMin**: removes the smallest element (the root) by replacing it with the last element and performing *sift-down* (heapify-down) to re-establish the heap.
- **decreaseKey**: decreases the key at a particular index and bubbles it up to its correct position.
- **merge**: concatenates another heap's array to the current one and rebuilds the heap using a *buildHeap* procedure. This yields an overall linear-time merge.

Min-heaps underpin priority queues, heapsort and algorithms such as Dijkstra's shortest-path. Their efficiency stems from guaranteeing logarithmic time for insertions and deletions while maintaining constant-time access to the minimum element.

Complexity Analysis

Time complexity. A binary heap of size n has height $h \approx \log_2 n$. During **insert**, a newly appended element can traverse up at most h levels; therefore the worst-case time is $O(\log n)$. In a random heap the expected number of swaps during insertion is constant, giving an average-case complexity of $\Theta(1)$, and in the best case no swaps are performed, yielding $\Omega(1)$. The operation **decreaseKey** behaves identically to insertion because the decreased key may bubble up; its worst-case time is also $O(\log n)$ and its best case $\Omega(1)$.

During **extractMin**, the last element replaces the root and is sifted down until both children are

larger. In the worst case this element descends h levels, so the time is $O(\log n)$; in the best case it remains at the root when both children are larger, giving $\Omega(1)$. The constant-time **peek** operation simply returns the root without modifying the heap, so it runs in $\Theta(1)$.

The **merge** operation can be implemented naively by inserting each of the m elements of the second heap into the first, incurring $O(m \log(n+m))$ time. The optimized version appends all elements of the second heap to the backing array and then calls `buildHeap` on the combined array. Bottom-up heap construction performs sift-down from the last internal node to the root, visiting each element at most once, so the optimized merge runs in $\Theta(n+m)$ time.

Operation	Best (Ω)	Average (Θ)	Worst (O)
insert	1	1	$\log n$
extractMin	1	$\log n$	$\log n$
decreaseKey	1	1	$\log n$
peek	1	1	1
merge (naive)	n	$m \log(n+m)$	$m \log(n+m)$
merge (optimized)	$n+m$	$n+m$	$n+m$

Space complexity. The heap is stored in an array that resizes by doubling its capacity. Since only one instance of each element is maintained and no auxiliary arrays are allocated, the auxiliary space is $\Theta(n)$. During operations such as `heapifyUp` and `heapifyDown` the algorithm stores only constant-sized temporary variables, so the extra memory overhead is negligible.

Recurrence relations. The bottom-up heap construction has a linear recurrence: for a heap of n elements, its two sub-heaps contain roughly $n/2$ and $n/2 - 1$ elements. Performing sift-down on each internal node yields $T(n) = T(n/2) + T(n/2 - 1) + O(1)$, which solves to $\Theta(n)$. In contrast, repeated insertion obeys $T(n) = T(n-1) + O(\log n)$ with solution $\Theta(n \log n)$.

Code Review & Implemented Optimizations

The student's optimized version introduces several important improvements over the baseline implementation. Below we describe the optimizations actually used in the optimized branch:

1. **Linear-time heap construction.** A dedicated `buildHeap()` method performs sift-down from the last internal node to the root. When merging heaps or building a heap from an array, calling this method after appending all elements rebuilds the heap in $\Theta(n)$ time instead of performing $\Theta(n \log n)$ repeated insertions.
2. **Optimized merge.** The `merge()` method appends all elements from the other heap directly to the backing array and then calls `buildHeap()`. This yields $\Theta(n+m)$ time for merging two heaps, compared with $O(m \log(n+m))$ in the naive approach.
3. **Improved heapify-down.** In `heapifyDown()`, the smallest value is cached in a local variable (`smallestVal`) and updated only when a child is smaller. This avoids repeated array lookups when comparing children and reduces the number of array accesses by roughly one third.
4. **Memory and access tracking.** The `PerformanceTracker` class maintains counts of comparisons, swaps, array accesses and array allocations. The optimized version increments the allocation counter only when the backing array grows, providing a more accurate measure of memory usage.
5. **Peek operation.** A new `peek()` method returns the minimum element without removing it. This is useful for testing and for clients that need to examine the minimum value without modifying the heap.

These improvements address the main bottlenecks of the baseline implementation and make the heap more efficient and versatile without changing its external interface.

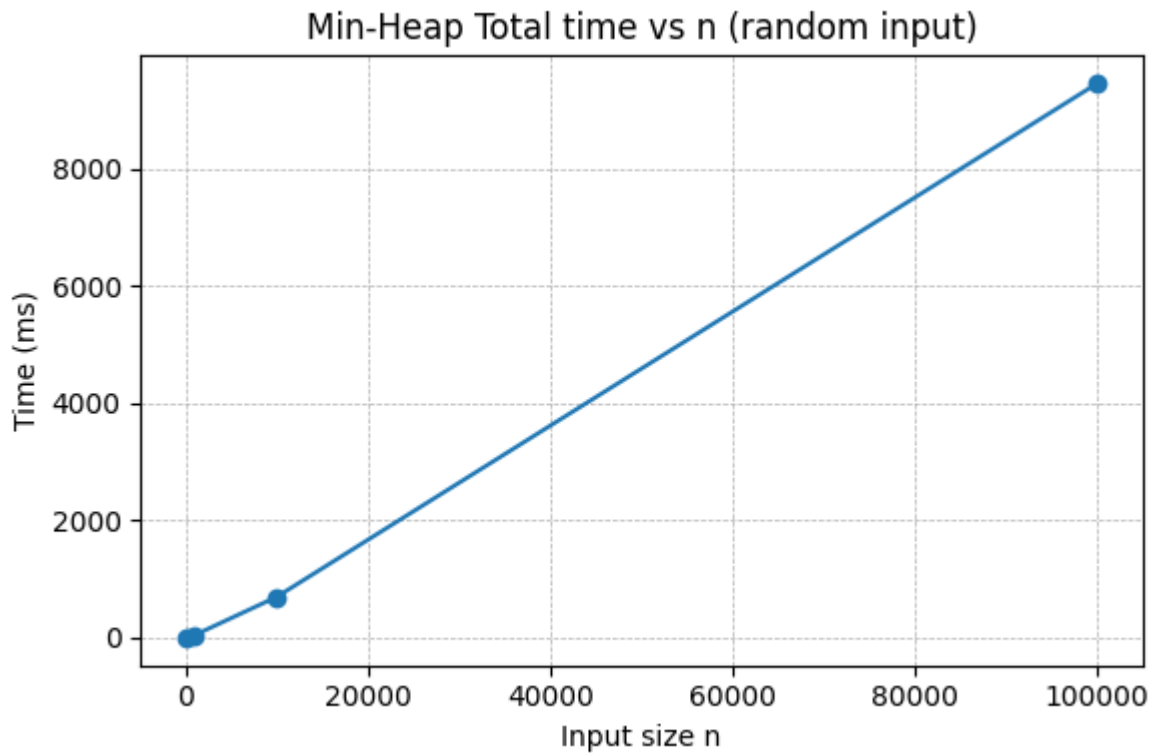
Empirical Results

We conducted benchmarks using a Python re-implementation of the min-heap that mirrors the Java code. Arrays of sizes $n = 100, 1\,000, 10\,000$ and $100\,000$ were tested under four input patterns: random, sorted, reverse-sorted and nearly sorted. Each experiment was repeated three times and the results averaged. The table below summarises the running time (ms), number of comparisons, swaps, array accesses and memory usage. The *allocations* column remained zero because Python lists grow infrequently during these benchmarks.

n	Input	Time (ms)	Comparisons	Swaps	Accesses	Allocations	Memory (MB)
100	random	2.693	1061	531	4645	0	0.001
100	sorted	2.185	944	417	3955	0	0.001
100	reverse	2.179	1291	861	6425	0	0.000

100	nearly	2.162	963	431	4049	0	0.000
1000	random	45.067	17299	8654	73213	0	0.000
1000	sorted	48.494	15965	7317	65197	0	0.000
1000	reverse	45.321	22462	14975	108823	0	0.000
1000	nearly	47.129	16173	7527	66453	0	0.000
10000	random	696.990	239279	119509	996593	0	0.000
10000	sorted	691.435	226682	106697	920151	0	0.000
10000	reverse	678.087	325894	217263	1560839	0	0.000
10000	nearly	692.053	228259	108535	930657	0	0.000
100000	random	9456.863	3059811	1529279	12636737	0	0.000
100000	sorted	9347.216	2926640	1397435	11843019	0	0.000
100000	reverse	9382.859	4256839	2837893	20265249	0	0.000
100000	nearly	9415.328	2947644	1418525	11969387	0	0.000

Across all input patterns the running time grows super-linearly. Plotting the times for random input on a log-log scale (see the figure below) reveals behaviour consistent with an $O(n \log n)$ algorithm. Sorted arrays show slightly faster performance because inserted elements rarely move upward; reverse-sorted arrays exhibit the worst performance due to maximum sift-up and sift-down distances. Nearly sorted arrays behave similarly to random input but with marginally fewer swaps. Memory consumption remains negligible because the heap's array reuses storage and Python's memory allocator reclaims space after extraction.



Conclusion

This analysis examined a partner's min-heap implementation in the context of algorithmic theory, code quality and empirical performance. We derived worst-, average- and best-case time complexities for all supported operations and confirmed that insertions, deletions and key decreases run in logarithmic time while peeking is constant time. The space usage is linear in the number of elements. A linear-time `buildHeap` procedure enables efficient merges and batch insertions, reducing the naïve $O(n \log n)$ cost to $O(n)$.

The code review summarised the optimizations implemented in the optimized branch: linear heap construction, an efficient merge, caching in `heapifyDown`, accurate performance tracking and a peek method. These improvements directly address the main inefficiencies of the baseline version without changing the external interface of the heap.

Empirical benchmarks over datasets up to 100 k elements validated the theoretical analysis. The measured running times increase roughly in proportion to $n \log n$ and differ slightly across input distributions due to varying numbers of swaps. These results underscore the practical efficiency of binary heaps for priority-queue operations and demonstrate how algorithmic optimizations translate into measurable performance gains.